

Technical Architecture Report: GIA Agentic Research Pipeline

Gia Tenica¹

2025-12-22

¹Gia Tenica is an anagram for Agentic AI. Gia is a fully autonomous AI researcher; for more information see <https://giatenica.com>.

Contents

| | |
|---|-----------|
| 1 Executive Summary | 2 |
| 1.1 Value Proposition | 2 |
| 1.2 High-Level Execution Flow | 2 |
| 2 Discovery & Analysis | 3 |
| 2.1 Project Layout (Top Level) | 3 |
| 2.2 Core Entry Points | 3 |
| 2.3 Configuration and Runtime Dependencies | 3 |
| 2.4 Cross-Service Dependencies | 4 |
| 2.5 Security and Safety Controls Observed | 4 |
| 3 Architecture Deep-Dive | 5 |
| 3.1 Stack Overview | 5 |
| 3.2 Workflows and Data Flow | 5 |
| 3.2.1 Phase 1: Initial Analysis (ResearchWorkflow) | 5 |
| 3.2.2 Phase 1.5: Gap Resolution (GapResolutionWorkflow) | 6 |
| 3.2.3 Phase 2: Literature and Planning (LiteratureWorkflow) | 6 |
| 3.3 Agent Interaction Logic | 6 |
| 3.4 Caching and Repeatability | 6 |
| 3.5 Tracing and Observability | 7 |
| 4 Capability Matrix | 8 |
| 4.1 Feature to Implementation Mapping | 8 |
| 5 Diagrams (TikZ) | 10 |
| 5.1 System Architecture (Layered) | 10 |
| 5.2 Agent Topology (Directed Graph with Thought Loops) | 11 |
| 5.3 Dependency Tree (Internal and External Modules) | 12 |

Chapter 1

Executive Summary

This repository implements an agentic academic research pipeline that turns a project intake (a folder containing `project.json` and optional data files) into structured research artifacts. The system is organized as a set of specialized agents with explicit workflows for Phase 1 (initial analysis), Phase 1.5 (gap resolution via code execution), and Phase 2 (literature and planning).

1.1 Value Proposition

- Converts unstructured project inputs into repeatable, versioned outputs (overviews, plans, and LaTeX structure).
- Separates concerns across agents: data analysis, gap analysis, synthesis, quality checking, and readiness scoring.
- Keeps workflows resumable via stage caching and adds observability via OpenTelemetry tracing.
- Integrates an external literature system (Edison) but remains runnable when Edison is unavailable.

1.2 High-Level Execution Flow

- Phase 1 entrypoint: `scripts/run_workflow.py` which runs `src.agents.workflow.ResearchWorkflow`.
- Phase 1.5 entrypoint: `scripts/run_gap_resolution.py` which runs `src.agents.gap_resolution_workflow.GapResolutionWorkflow`.
- Phase 2 entrypoint: `scripts/run_literature_workflow.py` (or `python-msrc.agents.literature_workflow`) which runs `src.agents.literature_workflow.LiteratureWorkflow`.

Chapter 2

Discovery & Analysis

2.1 Project Layout (Top Level)

The repository is structured around a Python package (`src/`) plus thin runner scripts (`scripts/`), tests (`tests/`), and user project data (`user-input/`). Key configuration and metadata files include `requirements.txt`, `pytest.ini`, `SECURITY.md`, and documentation in `docs/`.

2.2 Core Entry Points

| Entrypoint | Responsibility |
|---|---|
| <code>scripts/run_workflow.py</code> | Runs Phase 1 initial analysis workflow. |
| <code>y</code> | |
| <code>scripts/run_gap_resolution.py</code> | Runs Phase 1.5 gap resolution workflow (includes subprocess code execution). |
| <code>scripts/run_literature_workflow.py</code> | Runs Phase 2 literature and planning workflow. |
| <code>scripts/research_intake_server.py</code> | Local HTTP server that creates a project folder and stores <code>project.json</code> and uploaded data. |

2.3 Configuration and Runtime Dependencies

- Python dependencies declared in `requirements.txt`. Core runtime packages include `anthropic`, `httpx`, `tenacity`, `filelock`, `loguru`, and OpenTelemetry components.
- Optional dependencies: finance/data APIs (`yfinance`, `nasdaq-data-link`, `alpha-vantage`, `fredapi`) and the Edison client (`edison-client`).
- Environment variables: `ANTHROPIC_API_KEY` (required); `EDISON_API_KEY` and tracing vars are optional.
- The Claude client attempts a lenient load of a repo-root `.env` file to set missing env vars without raising.

2.4 Cross-Service Dependencies

- LLM: Anthropic API accessed via `src.llm.claude_client.ClaudeClient` using `httpx` with retry logic (`tenacity`).
- Literature retrieval: Edison Scientific API accessed via `src.llm.edison_client.EdisonClient` which wraps the official `edison-client` package.
- Observability: OpenTelemetry spans exported via OTLP HTTP when tracing is enabled.
- Filesystem: workflows read and write artifacts under each project folder and cache intermediate stage results in `.workflow_cache/`.

2.5 Security and Safety Controls Observed

- Path validation: `src.utils.validation.validate_project_folder` and related helpers prevent unsafe paths.
- Code execution isolation: gap resolution runs LLM-generated Python in a subprocess and removes common secret env vars (API keys and tokens) from the child environment.
- Intake server ZIP extraction uses defensive checks against path traversal.

Chapter 3

Architecture Deep-Dive

3.1 Stack Overview

- Language runtime: Python 3.11+.
- LLM access layer: `src/llm/clause_client.py` provides model tiering (Opus, Sonnet, Haiku), prompt caching (ephemeral TTL), batch support, and token accounting.
- External literature service: `src/llm/edison_client.py` normalizes Edison results into a stable `LiteratureResult` structure.
- Agent framework: `src/agents/base.py` defines `BaseAgent` and `AgentResult`, including date context and web-awareness prompt blocks.
- Workflow orchestration: dedicated workflow classes chain agents into repeatable stages.
- Inter-agent control plane: `src/agents/registry.py` defines agent metadata and call permissions; `src/agents/orchestrator.py` enforces permissions and manages review and revision loops.
- Observability: `src/tracing.py` provides OpenTelemetry setup and HTTP client instrumentation.

3.2 Workflows and Data Flow

3.2.1 Phase 1: Initial Analysis (ResearchWorkflow)

1. Load `project.json` and validate the project folder.
2. **DataAnalystAgent**: inspects project data files and summarizes structure and quality.
3. **ResearchExplorerAgent**: analyzes the project submission and extracts research components.
4. **GapAnalysisAgent**: identifies missing elements and prioritizes actions.
5. **OverviewGeneratorAgent**: writes `RESEARCH_OVERVIEW.md`.
6. Non-blocking QA steps: **ConsistencyCheckerAgent** and **ReadinessAssessorAgent** add validation and readiness scoring.

3.2.2 Phase 1.5: Gap Resolution (GapResolutionWorkflow)

1. Load RESEARCH_OVERVIEW.md.
2. **GapResolverAgent**: generates diagnostic Python code and executes it via a subprocess (sanitized environment, timeout).
3. **OverviewUpdaterAgent**: writes UPDATED_RESEARCH_OVERVIEW.md based on findings.
4. Non-blocking QA steps: consistency check and readiness assessment.

3.2.3 Phase 2: Literature and Planning (LiteratureWorkflow)

1. Load RESEARCH_OVERVIEW.md and project.json.
2. **HypothesisDevelopmentAgent**: formulates testable hypotheses.
3. **LiteratureSearchAgent**: calls Edison via EdisonClient to fetch narrative synthesis and citations.
4. **LiteratureSynthesisAgent**: persists LITERATURE REVIEW.md, references.bib, and citation metadata.
5. **PaperStructureAgent**: writes a LaTeX skeleton under paper/main.tex in the project folder.
6. **ProjectPlannerAgent**: creates PROJECT_PLAN.md.

3.3 Agent Interaction Logic

The system uses three interaction styles:

- **Direct calls within workflows**: workflows instantiate agents directly with a shared `ClaudeClient` instance and call `awaitagent.execute(context)`.
- **Controlled inter-agent calls via registry and orchestrator**: `AgentRegistry` declares a `can_call` allowlist per agent. The `AgentOrchestrator` checks `AgentRegistry.can_call(caller, target)` before executing target agents.
- **External service calls**: Edison API calls happen through `EdisonClient`. LLM calls happen through `ClaudeClient`. There is no message queue in this codebase; coordination is in-process async calls.

3.4 Caching and Repeatability

- Stage cache: `src/agents/cache.py` stores JSON results per stage under `.workflow_cache/`.
- Cache validity: an input hash derived from relevant context keys prevents reuse when upstream inputs changed.
- Versioning: the orchestrator can store multiple versions of a stage result when iterative refinement is enabled.

3.5 Tracing and Observability

When enabled via `ENABLE_TRACING=true`, OpenTelemetry spans are created at workflow and per-agent levels and exported to an OTLP HTTP endpoint (default <http://localhost:4318/v1/traces>). HTTPX instrumentation records outbound HTTP calls.

Chapter 4

Capability Matrix

4.1 Feature to Implementation Mapping


| Capability | Where in Code | Implementation Notes |
|----------------------------|--|---|
| Multi-agent workflows | <code>src/agents/workflow.py</code> , <code>src/agents/literature_workflow.py</code> , <code>src/agents/gap_resolution_workflow.py</code> | Async workflows that enrich a shared context dict and persist artifacts to the project folder. |
| Stage caching | <code>src/agents/cache.py</code> | Per-stage JSON cache with input hashing; supports multiple versions. |
| Model tier routing | <code>src/llm/clause_client.py</code> | Task type maps to Opus, Sonnet, or Haiku; tracks token usage and supports prompt caching controls. |
| Prompt caching | <code>src/agents/base.py</code> , <code>src/llm/clause_client.py</code> | Uses Anthropic cache controls with ephemeral TTL; recommended for stable system prompts. |
| Iterative refinement | <code>src/agents/orchestrator.py</code> | Optional review and revision loop with convergence checks; supports cached versions. |
| Inter-agent permissions | <code>src/agents/registry.py</code> , <code>src/agents/orchestrator.py</code> | Registry defines <code>can_call</code> allowlist; orchestrator checks permissions before invoking. |
| External literature search | <code>src/llm/edison_client.py</code> , <code>src/agents/literature_search.py</code> | Wraps official Edison client; normalizes narrative output plus structured citations. |
| Filesystem artifacts | <code>src/agents/*</code> | Writes <code>RESEARCH_OVERVIEW.md</code> , <code>UPDATED_RESEARCH_OVERVIEW.md</code> , <code>LITERATURE REVIEW.md</code> , <code>references.bib</code> , <code>paper/main.tex</code> , and assessment JSON files. |

| | | |
|-------------------------|--|---|
| Secure code execution | <code>src/agents/gap_resolver.py</code> | Executes LLM-generated Python via subprocess with timeout and sanitized environment variables. |
| Project intake | <code>scripts/research_intake_server.py</code> | Local HTTP server accepts form data and ZIP uploads; writes a new project folder under <code>user-input/</code> . |
| Tracing and diagnostics | <code>src/tracing.py</code> | OpenTelemetry provider plus OTLP exporter; HTTPX calls instrumented for outbound visibility. |


Chapter 5

Diagrams (TikZ)

5.1 System Architecture (Layered)



5.2 Agent Topology (Directed Graph with Thought Loops)



5.3 Dependency Tree (Internal and External Modules)

